

gap-closure: TF-IDF -156%    Qwen2.5-0.5B -94%    Qwen2.5-7B (chat) -44%    (all 8 architectures UNDER cmaj baseline)

Verifier-encoder scaling: chat-LMs monotone, but embedding/math tunings hurt

